

ELEC 377

Operating Systems

Week 4 Lecture 1

Review

Cannot cover the three weeks in depth, just a quick overview, not comprehensive. There may be material on the quiz not covered today.

Reminder the first Quiz is Tomorrow
- covers weeks 4-6 and lab 2

First Come First Served (FCFS)

- Simple, easy to implement
 - ◊ Ready queue is a first-in-first-out (FIFO) queue
 - ◊ Non-preemptive. Once in the Running state, the process stays running until it asks for I/O or exits.
- Average Waiting time may be long
 - ◊ Short Bursty I/O do not have priority over CPU intensive jobs
 - ◊ variance in wait time / throughput is large, depends on order of jobs
 - ◊ *Convoy effect*, all I/O jobs end up behind CPU jobs which hog the CPU
- Tends to make poor response in interactive systems
 - ◊ used in simple OS or in systems with very little variance in CPU burst times

Shortest Job First (SJF)

- scheduling ordered on the length of the next CPU burst
- Nonpreemptive - process gets entire burst time like FCFS

Estimating CPU Burst Times - Exponential Average

- Use length of last CPU burst to predict next burst

t_n = current CPU burst time

τ_0 = initial estimate

τ_n = predicted for current burst

τ_{n+1} = prediction for next CPU burst

α = weighting parameter

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

$\alpha = 0 \rightarrow \tau_{n+1} = \tau_0$ (initial estimate) never changes

$\alpha = 1 \rightarrow \tau_{n+1} = t_n$ (last time slice) only used

Estimating CPU Burst Times

- predicted time always lags real time
- If process spends a reasonable period of time at a constant burst range then estimate approaches current burst time
- what is reasonable? how to tune?
 - ◇ α is the tuning parameter
 - ◇ α is low, then past behaviour has more weight, the estimate is slower to change
 - ignore transient behaviour
 - ◇ α is high, then last time slice has heavier weight, the estimate is faster to change
 - more susceptible to transient behaviour

Shortest Remaining Time First (SRTF)

- SJF was non-preemptive
 - ◇ process burst gets entire time needed
- Preemptive version
 - ◇ after an interrupt, calculate remaining time for current burst
 - ◇ will be less than all other process in ready queue
 - ◇ if a new process becomes ready with shorter burst than remaining
 - dispatch the new process

SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

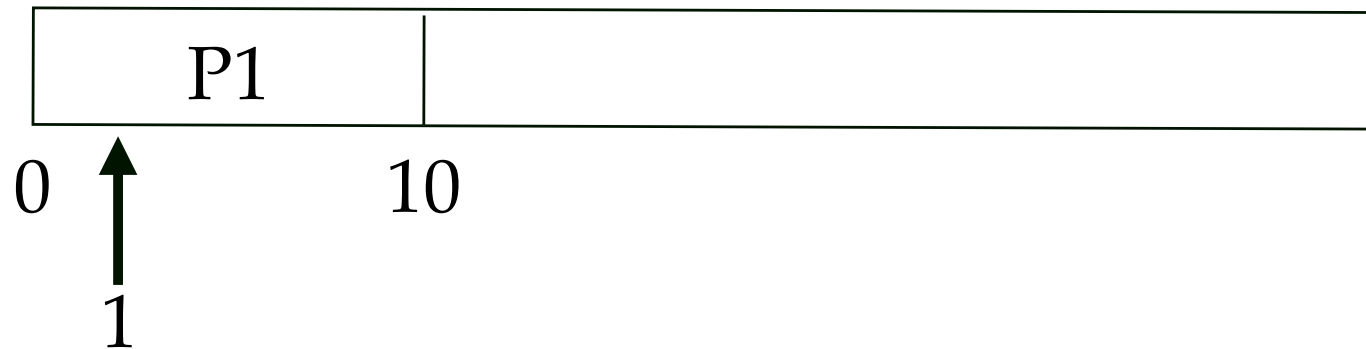
| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |



0

SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |



| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

P2 5

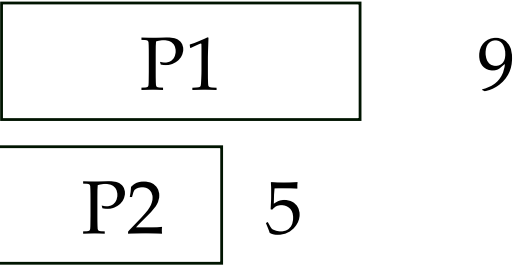
SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |



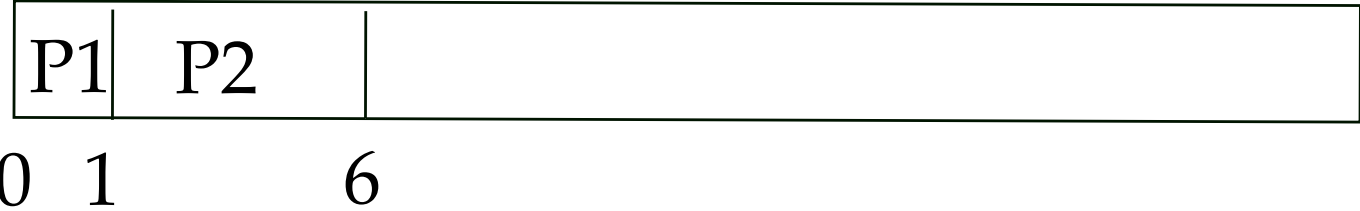
0 1

| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |



SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

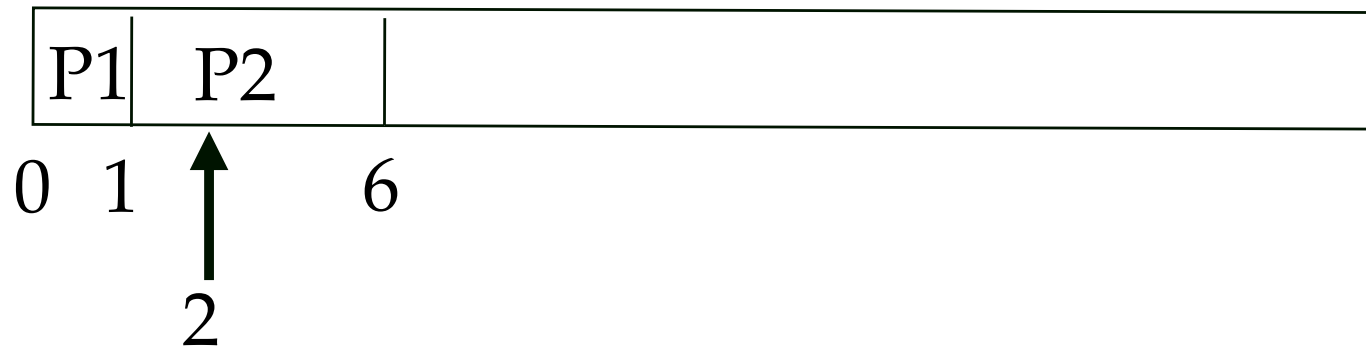


| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

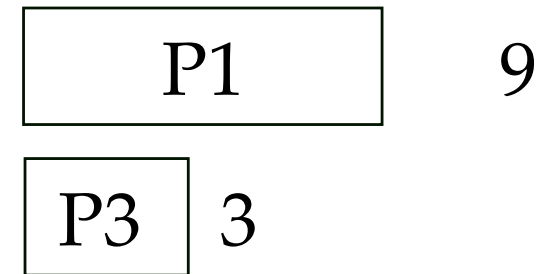


SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

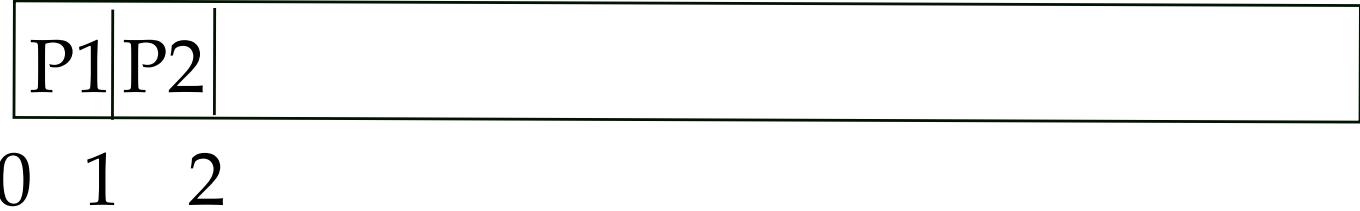


| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

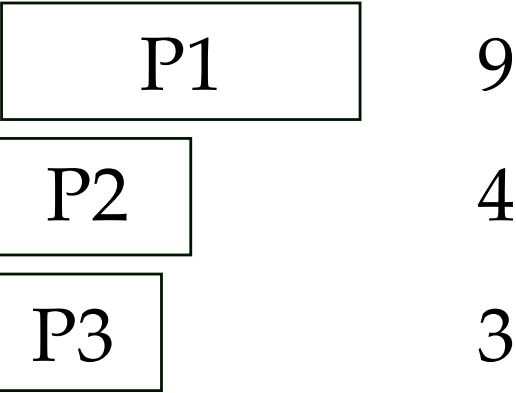


SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

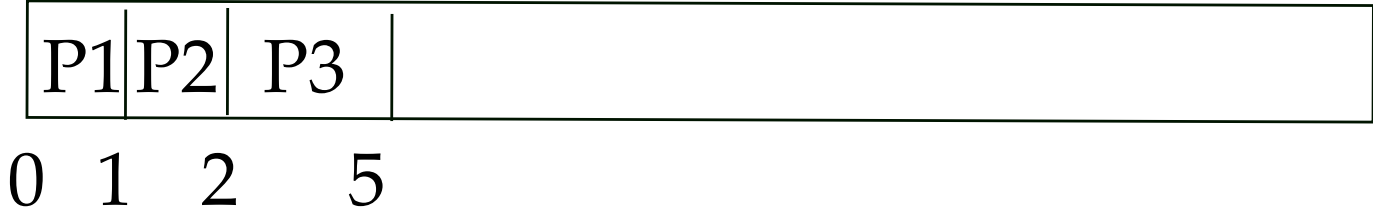


| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

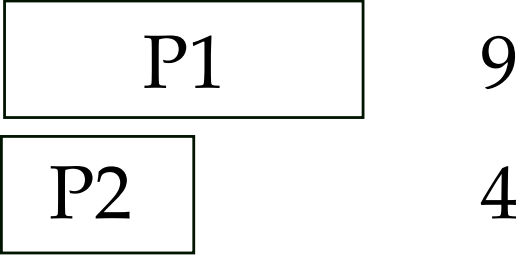


SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

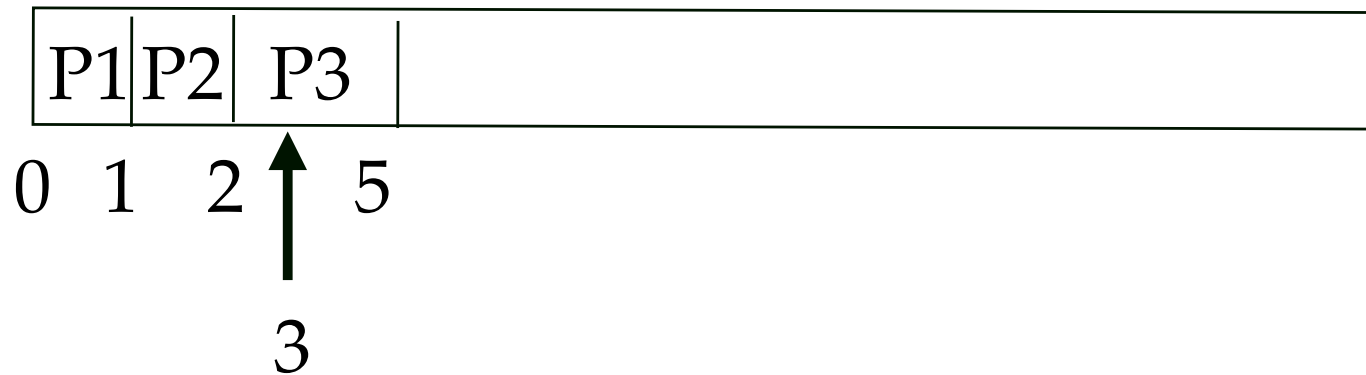


| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

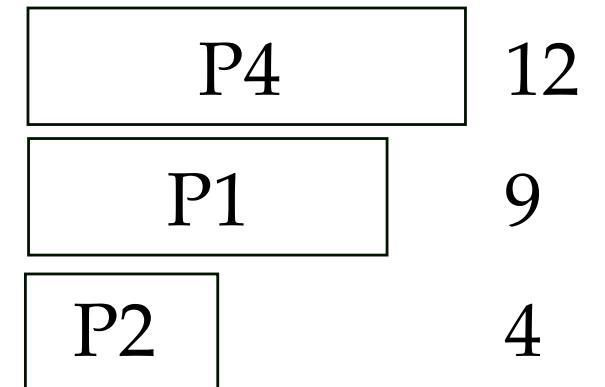


SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

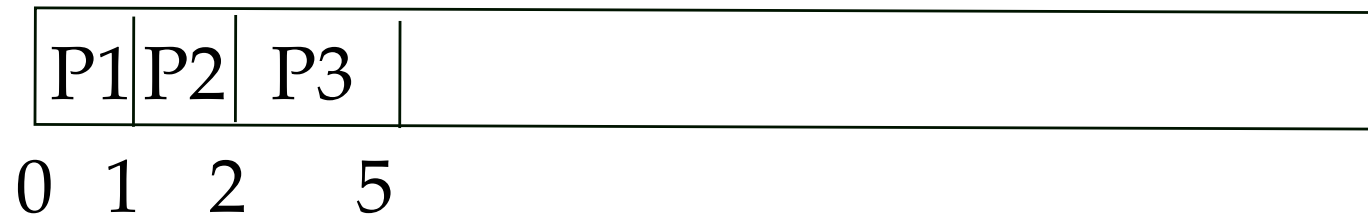


| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

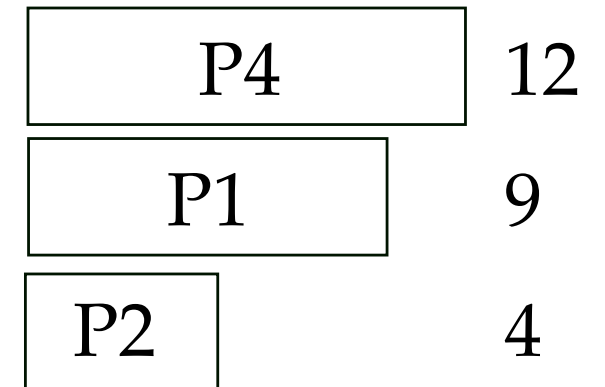


SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |

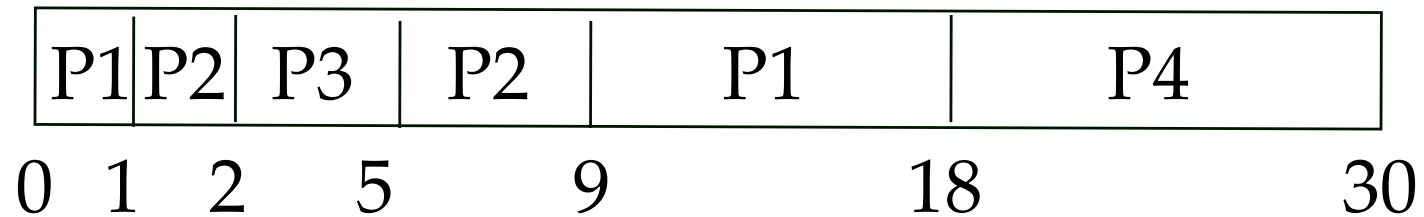


| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |



SRTF Example

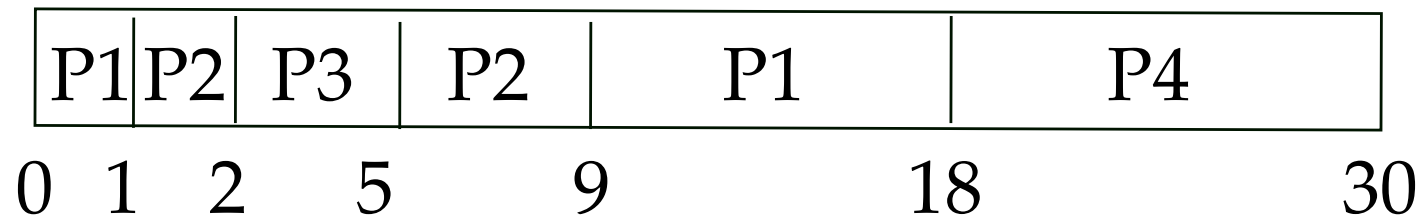
| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |



| # | Expr | Wait |
|---|------|------|
| | | |
| | | |
| | | |
| | | |

SRTF Example

| # | Arrival | Length |
|----|---------|--------|
| P1 | 0 | 10 |
| P2 | 1 | 5 |
| P3 | 2 | 3 |
| P4 | 3 | 12 |



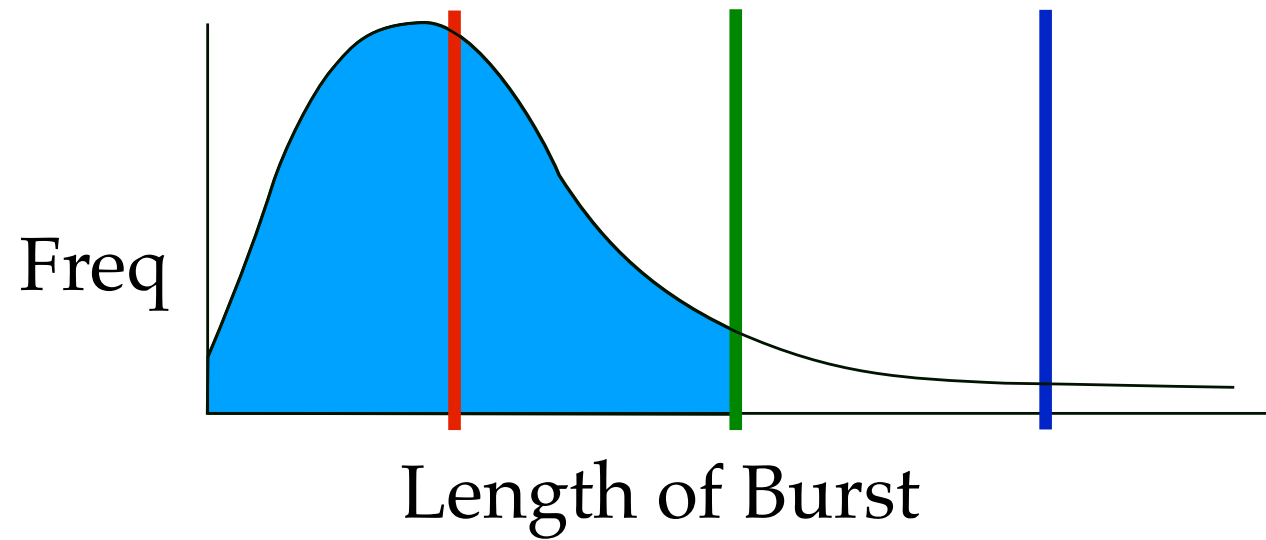
| # | Expr | Wait |
|----|-------------------|------|
| P1 | $0 + (9 - 1)$ | 8 |
| P2 | $(1-1) + (5 - 2)$ | 3 |
| P3 | $2 - 2$ | 0 |
| P4 | $18 - 3$ | 15 |

- Total Wait = 26 ms
- Average Wait = 6.5 ms

Round Robin (RR)

- Similar to FCFS, but with preemption.
- Designed for Time Sharing Systems
- Time slice (*quantum*) → maximum time process gets to run
- If the quantum (q) is large → FCFS
- If q is small, then appears to be multiple slower CPU's.
- Context switching is not free
 - ◇ shorter q, more context switches to complete a single CPU burst for a given process
 - ◇ q must be large with respect to context switch time

Round Robin – Quantum Length



- 80% of CPU bursts should be shorter than q .

Priority Scheduling

- Each process has a priority
- CPU goes to process with highest priority
 - ◇ ready queue is sorted based on priority
 - ◇ process with the same priority is handled FCFS
 - ◇ preemptive / nonpreemptive
 - ◇ internal/external
 - ◇ starvation

Multilevel Queues

highest priority

System Processes

Interactive Processes

Interactive Editing Processes

Batch Processes

Experimental Processes

lowest priority

Multi Level Queues

- Each queue has it's own scheduling algorithm
- Interactive (foreground) - probably Round Robin
- Scheduling must be done between the queues
 - ◊ usually fixed priority preemptive scheduling (starvation)
 - ◊ time slice between queues (portion time between queues)
- In simplest form, processes are assigned a queue and remain there until completion
- Higher priority queues may require more money, or more status

Multi Level Feedback

- processes move between queues
- when doing I/O, processes move to higher priority queues
- When CPU intensive, processes move to lower priority queues
- Give higher priority queues smaller quanta (preemptive)
- Processes that use entire quanta are too high priority, bump down to lower priority queue
- Processes that don't use entire quanta are too low priority and moved up to a higher priority queue

Multi Level Feedback

- parameters
 - ◇ number of queues
 - ◇ the scheduling algorithm for each queue
 - ◇ when to upgrade a process
 - ◇ when to downgrade a process
 - ◇ how to choose the initial queue
- most complex algorithm, is approximated using priorities

System Model

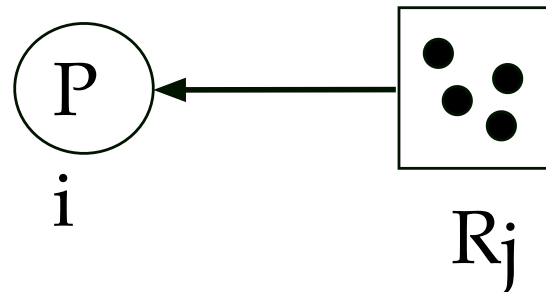
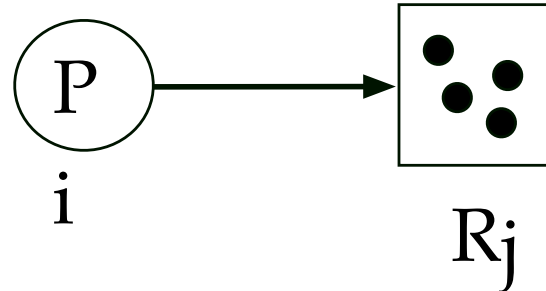
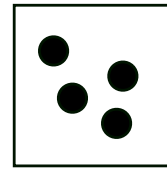
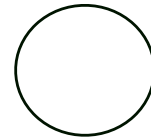
- Resource Types R_1, R_2, \dots, R_n
 - ◇ Each resource has a number of instances (W_i)
 - ◇ Resource instances are indistinguishable
 - doesn't matter which one you get.
- Process resource protocol
 - ◇ request
 - ◇ use
 - ◇ release

Deadlock Conditions

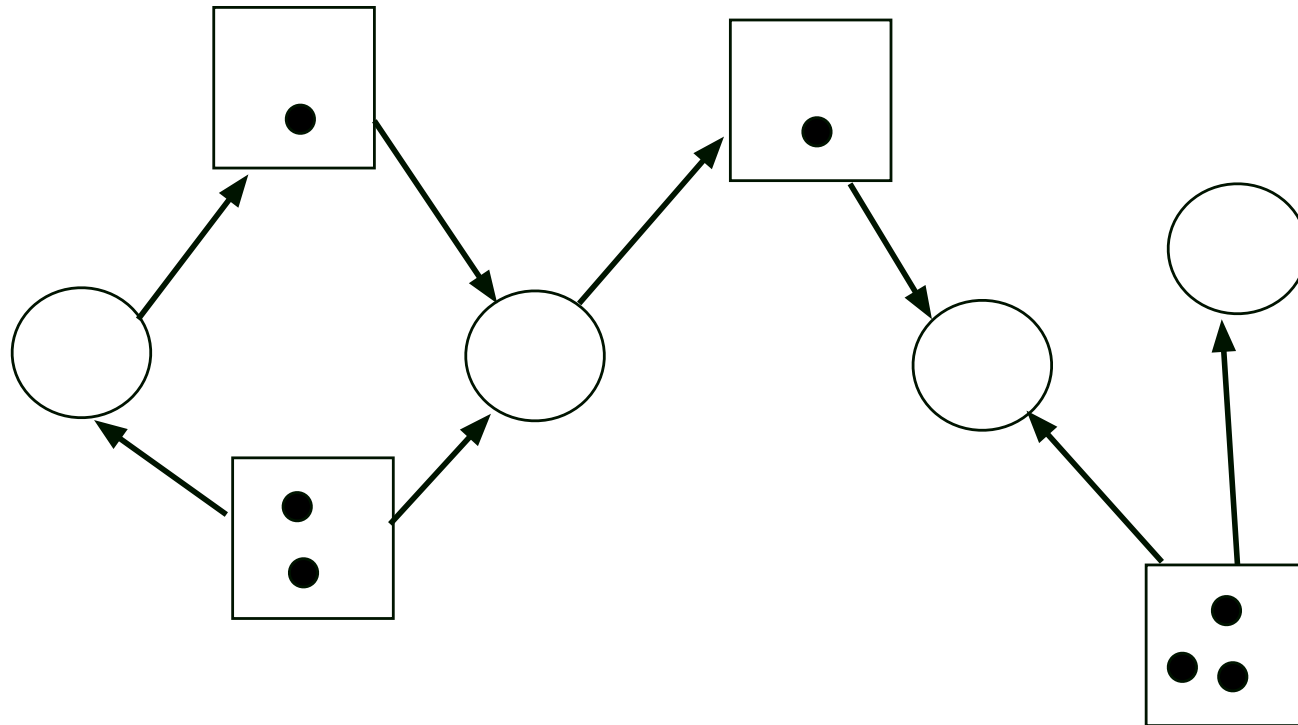
- four conditions necessary for deadlock:
 - ◇ **mutual exclusion:** only a limited number (usually one) process at a time can use a resource
 - ◇ **hold and wait:** a process has (at least) one resource and is waiting for another
 - ◇ **no preemption:** we can't take a resource away from a process
 - ◇ **circular wait:** P_0 waits for a resource held by P_1 , which waits for a resource held by $P_2, \dots P_n$, which waits for a resource held by P_0

Resource Allocation Graph

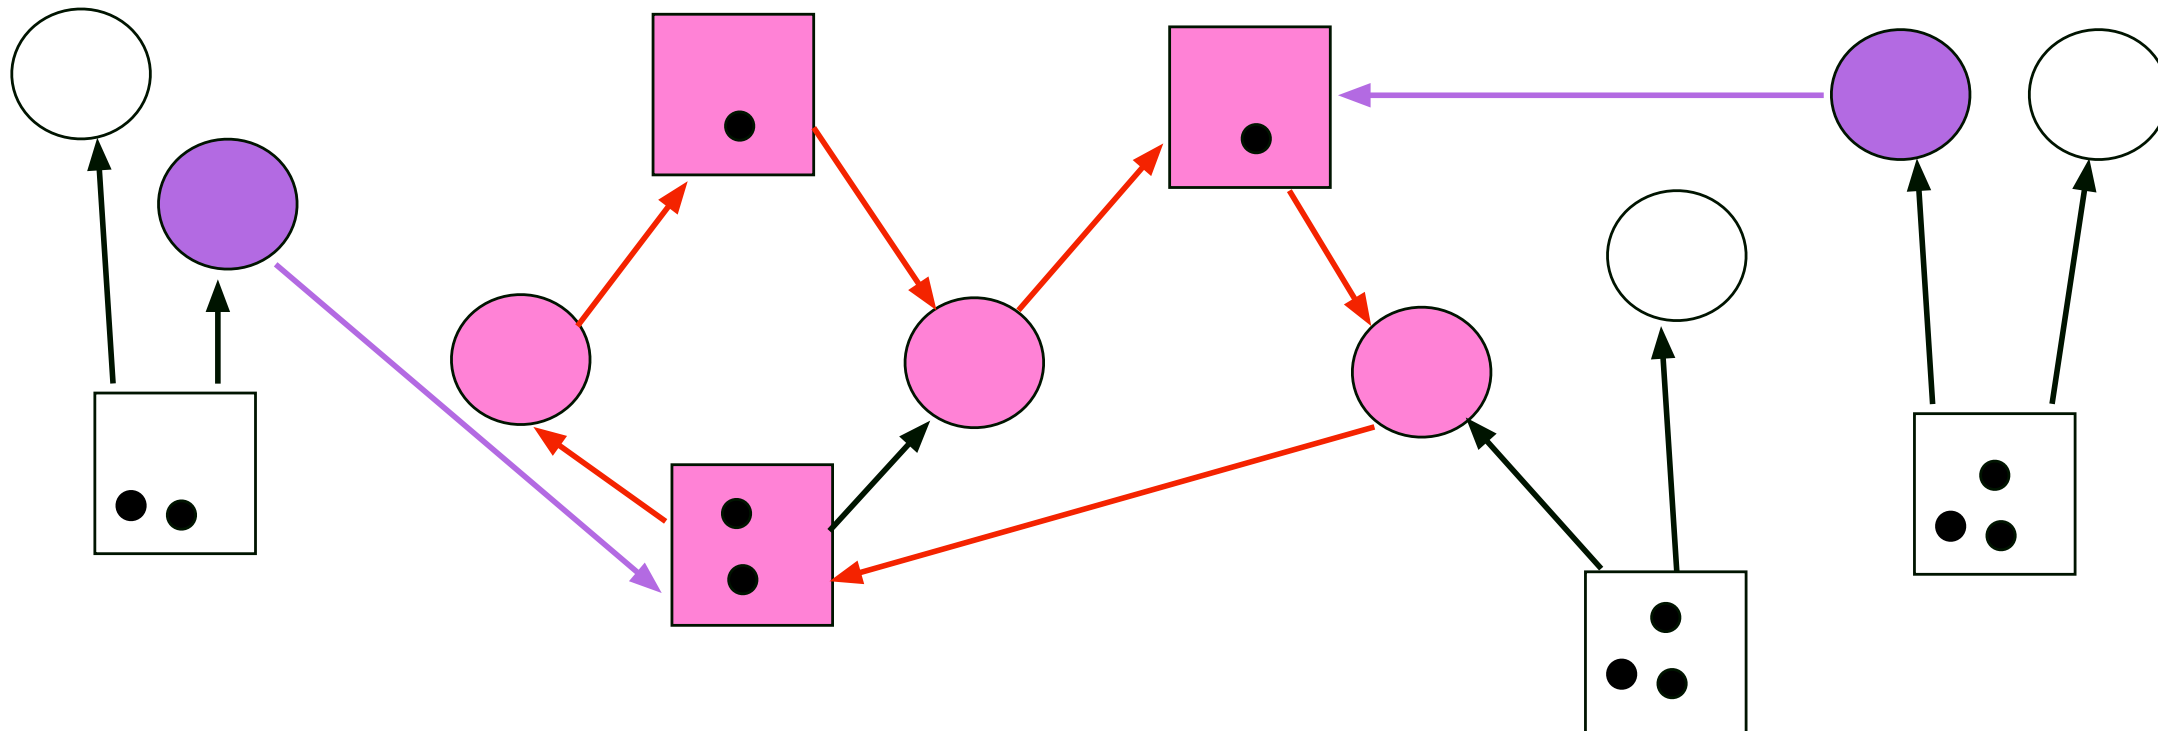
- Process
- Resource Type
◊ 4 instances
- P_i requests an instance of R_j
- P_i holds an instance of R_j



Resource Allocation Graph Example



Deadlock Example



What do we do??

- Prevention
 - ensure one of the 4 conditions never happens
- Avoidance:
 - extra information before allocating an available resource
- Recovery:
 - enter deadlock state and recover
- Ignore
 - hope it never happens
 - handle it manually
 - most interactive operating systems use this approach

Safe State

- system is safe if there is some order we can allocate the resources and not produce a deadlock
 - ◇ might not be the order that the processes actually request the resources
 - ◇ safe order means that one of the processes may have to wait
- $\langle P_1, P_2, \dots, P_n \rangle$ is safe if P_i can satisfy the maximum resources with available and the resources owned by previous processes.
 - ◇ P_1 max must be satisfied only with free resources
 - ◇ P_2 max must be satisfied with free + P_1
 - ◇ P_3 max gets available + P_1 + P_2
 - ◇ If not, wait until a previous process finishes.

Safe State – Examples

| Process | Current | Max |
|---------|---------|-----|
| P0 | 5 | 10 |
| P1 | 2 | 4 |
| P2 | 2 | 9 |

Total = 12, Free = 3

$\langle P1, P0, P2 \rangle$

$P1(2) + 3 = 5 \geq P1Max(4)$

$5 + P0(5) = 10 \geq P0Max(10)$

$5 + P2(2) = 7 \text{ not } \geq P2Max(9)$

$10 + P2(2) = 12 \geq P2Max(9)$

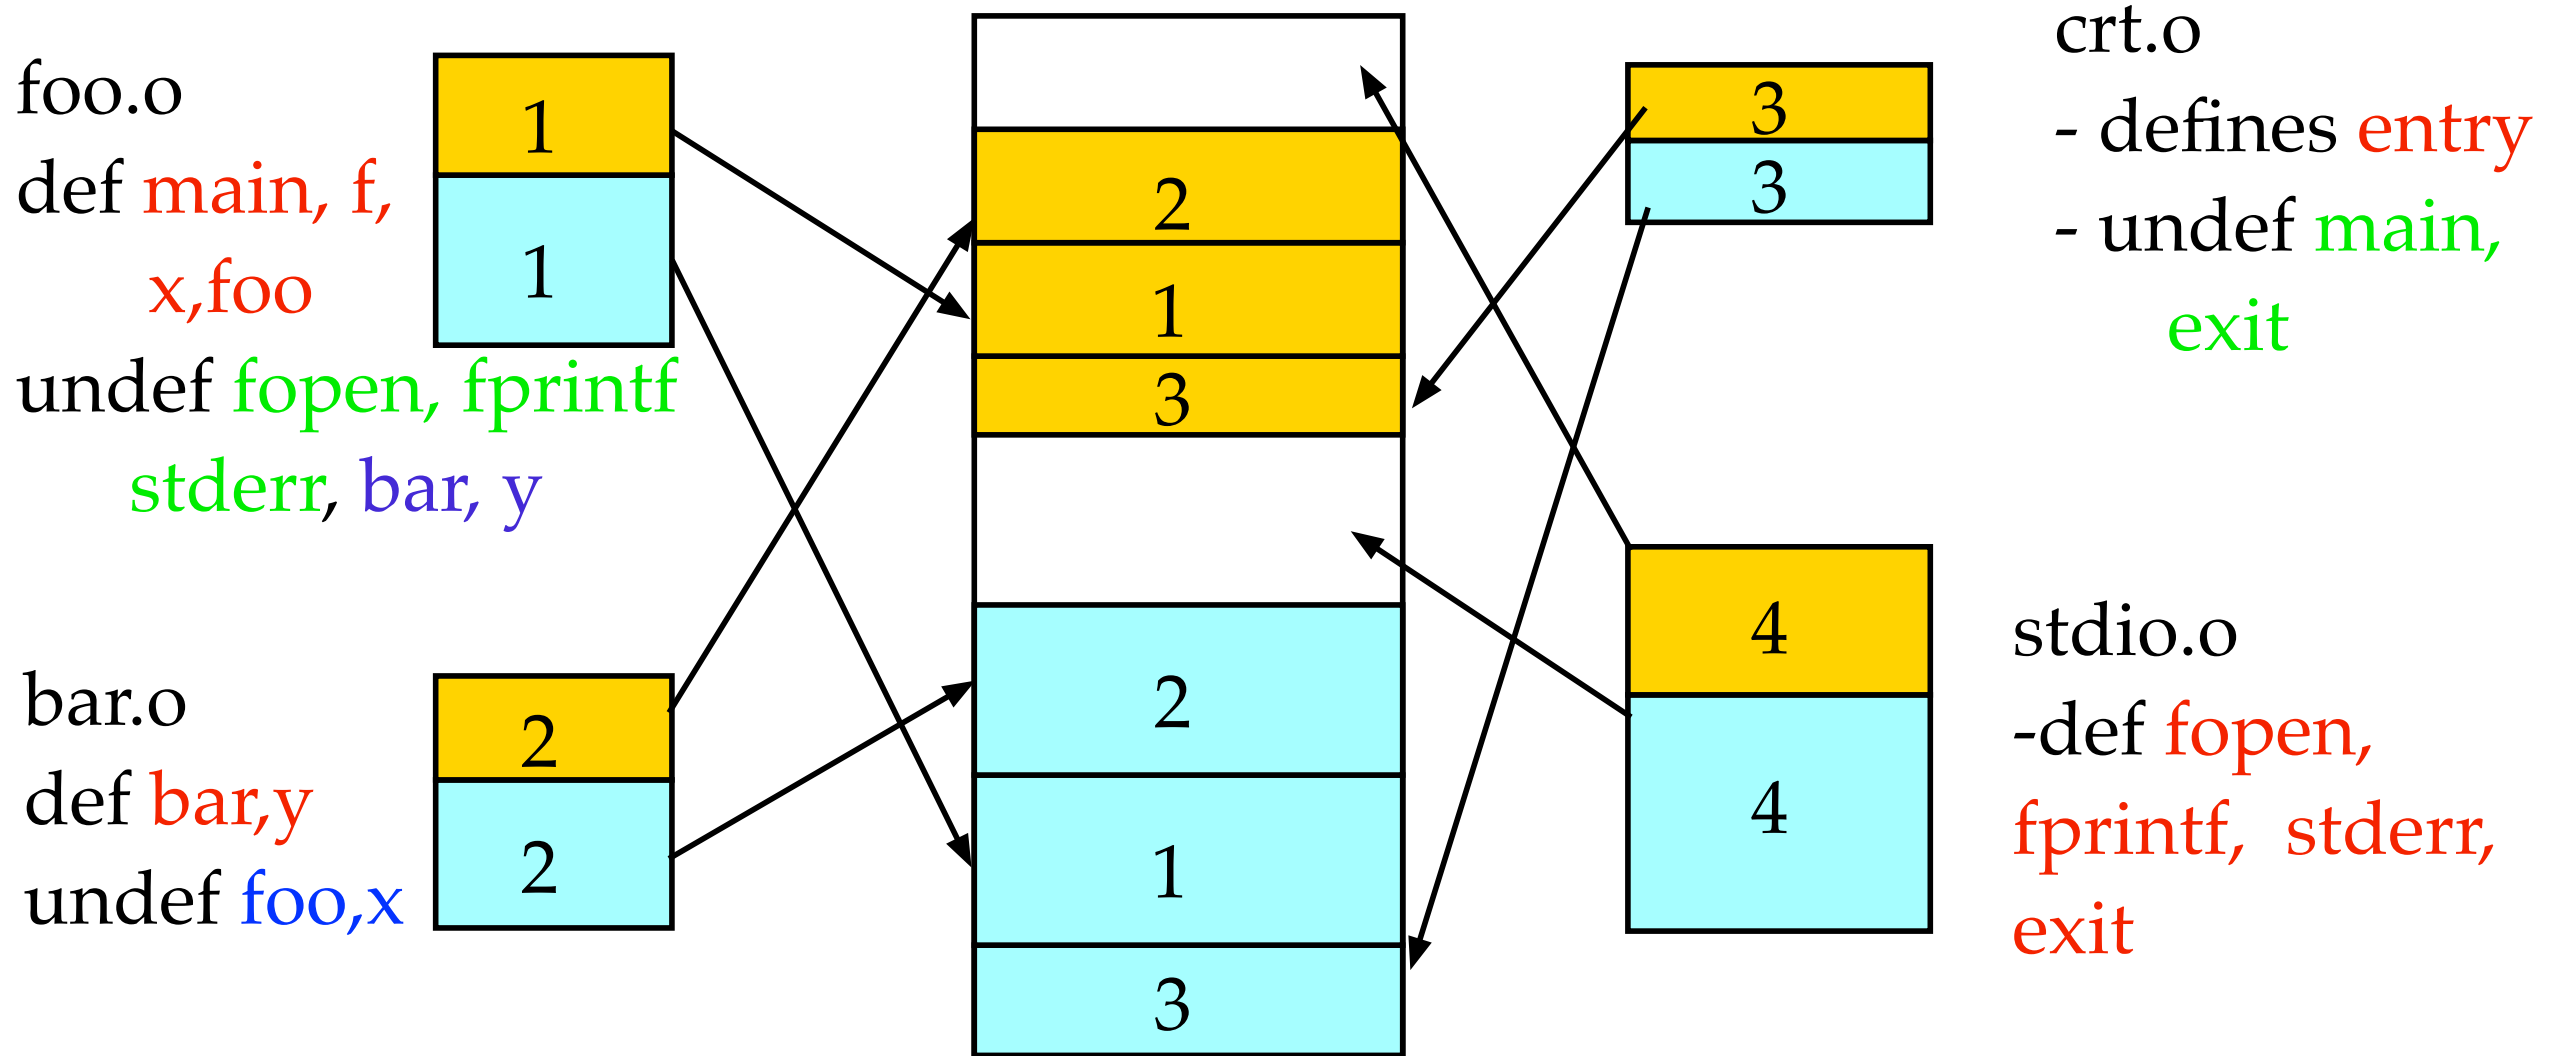
Binding Instructions and Data

- Entities in the original program must be bound to a location in memory
 - ◊ `int count;`
- Programs may reside in different parts of memory
- Three different stages
 - ◊ Compile (and linkage) time (MS-DOS .COM, Arduino)
 - ◊ Load Time – When loader loads program into memory, addresses are resolved (early Unix, Linux)
 - ◊ Execution Time – programs can move in memory
 - hardware support required

Compilation Process

- This presentation describes how several separate programs are compiled and assembled (linked) into an executable.

Linking—combine to single executable



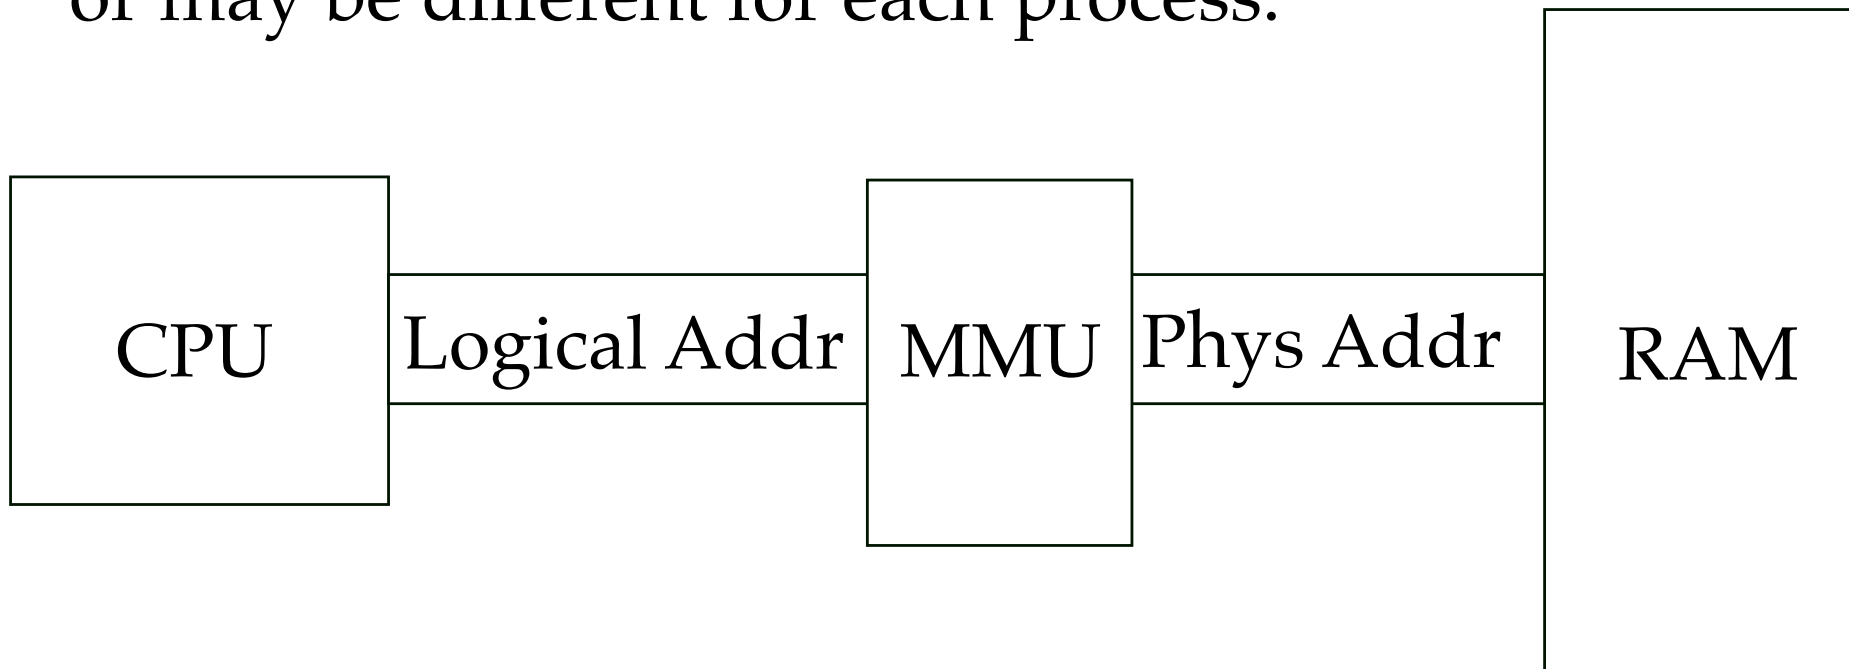
Logical vs Physical Address Space

Central Concept to Memory Management

- Logical Address
 - ◇ address generated by CPU
 - ◇ also known as a virtual address
- Physical Address
 - ◇ location in physical memory
- Logical and Physical address are the same in compile and load time address binding. They differ in execution time binding
- User program only deals with logical addresses. It never sees the physical address

Memory Management Unit (MMU)

- Hardware that maps virtual to physical address
 - ◇ many different approaches
- Logical address may be shared between process, or may be different for each process.



Contiguous Memory

- Using relocation and limit register assumes that the process is given a contiguous chunk of physical memory.
 - ◊ single partition allocation
 - ◊ simple allocation strategy, similar to that used by malloc in C or new in Java.
- main memory is divided into two parts
 - ◊ operating system (usually in same part as the interrupt vector)
 - ◊ User memory (divided among processes)

Storage Allocation

Three general approaches

- First Fit
 - ◇ use the first hole on the free list that is big enough
 - ◇ only look at part of list
- Best Fit
 - ◇ smallest block that is large enough
 - ◇ search entire list
- Worst Fit
 - ◇ largest block (largest remainder)
 - ◇ worst algorithm

Fragmentation

- Internal Fragmentation
 - ◇ if we allocate memory in units larger than a single byte (say 1K)
 - ◇ last block is only partially used
- External Fragmentation
 - ◇ lots of small holes spread throughout memory, none are big enough to satisfy a request
 - ◇ worst fit tries to reduce this
 - ◇ compaction - move blocks (requires execution-time binding)

Shared Libraries

- *Static* linking is when the all of the modules including system libraries are linked together at compile time.
 - ◇ shared libraries must be preallocated into common address space at runtime.
- When dynamic linking, the OS invokes the dynamic linker when the process is loaded.
 - ◇ linking uses the Program Link Table (PLT)
 - ◇ The code must be position independent.

Paging

- Why should memory have to be contiguous?
- Physical memory is divided into frames
 - ◊ typically 512 bytes to 8K in size
 - ◊ Linux is 4K
- Logical memory is divided into pages
 - ◊ same size as frames
- If process needs n pages, find n free frames in memory
 - ◊ no need to be contiguous
- A table translates from each page to the appropriate frame
- No external fragmentation, but still have internal fragmentation

Paging

- Logical Address Space and Physical Address space may not be the same size!!
 - ◊ physical address may be larger
(e.g. 32 bit logical, 40 bit physical)
 - ◊ physical address may be smaller
(64 bit logical = $1.8e19$ bytes)
- Frame and page always the same size
 - always power of 2

Paging - TLB

- Want to minimize extra memory traffic of page tables
- Small cache inside of MMU
 - ◇ TLB - translation look-aside buffer
 - ◇ associative memory

| Page # | Frame # |
|--------|---------|
| | |
| | |
| | |
| | |
| | |

Sharing Pages

- Shared Libraries
- Multiple invocations of a given program (e.g. shell, editor)
- Contiguous memory allocation made sharing difficult
- Shared code must be reentrant
 - ◊ Must not modify itself
 - ◊ Must also be position independent
- Most data is not shared
 - ◊ however IPC Shared Segments are now easy!!
- Page table entries for shared code and data in each process point to the common frames
- Page table entries for private data point to different frames

Page Structures - Summary

- Three ways of reducing the memory requirements of the page tables
- *All* of them increase the cost of converting a logical address to a physical address
 - ◇ TLB absorbs much of the cost
 - ◇ Increase the cost of a TLB miss
 - ◇ Effectiveness of TLB is more important
- Hardware vs software table walk

Virtual Memory

- Not all of the program need be in memory at one time
 - ◇ dynamic loading, overlays
- Some code may never be executed for a given instance of a program.
 - ◇ A word processor may have to store inline video in a document.
- Not all of the code or data in a program is used at the same time
 - ◇ compilers read and analyze source code before generating output assembly code
 - ◇ word processors only read from a file on an open command.
- If the code is not executing, why should it be in memory?